

Muhammad Anas

Senior AI Engineer

Location: Lahore, Pakistan | **Phone:** +92 303 7884884 | **Email:** anas.dev72@gmail.com

LinkedIn: <https://www.linkedin.com/in/anas72/>

Senior AI Engineer with 5+ years of experience in designing and deploying production-grade AI systems, scalable backend architectures, and LLM-powered applications. Specialized in RAG pipelines, AI agents, embeddings-based retrieval, and LLM fine-tuning (RLHF). Experienced in building and scaling AI microservices using Python (FastAPI), Node.js, React.js, and Next.js, deployed on AWS, GCP, and Azure using Docker and Kubernetes. Strong expertise in AI automation workflows (n8n), vector databases (Pinecone, FAISS), and event-driven architectures, delivering measurable improvements in performance, scalability, and automation efficiency.

Technical Skills

- **Programming Languages:**
Python, JavaScript, TypeScript
- **AI / Machine Learning:**
LLMs, RAG, Prompt Engineering, LangChain, LangGraph, LlamaIndex
Transformers, RLHF, Fine-Tuning, AI Agents, NLP
Multi-Agent Systems, Tool Calling, Function Calling
- **LLM & AI Ecosystem:**
OpenAI, Hugging Face, Claude, Gemini APIs
Embeddings, Semantic Search, Vector Search Optimization
- **Backend Development:**
FastAPI, Django, Node.js, REST APIs, GraphQL
Microservices Architecture, Event-Driven Systems
- **Frontend Development:**
React.js, Next.js, JavaScript (ES6+)
- **Cloud & DevOps:**
AWS, GCP, Azure
Docker, Kubernetes, Serverless (Lambda)
CI/CD Pipelines, GitHub Actions
- **Automation & AI Workflows:**
n8n, Zapier, Workflow Automation
AI Orchestration, Event-Based Systems
- **Databases:**
PostgreSQL, MongoDB, Redis
- **Tools:**
Git, GitHub, Agile/Scrum

Professional Experience

Senior AI Engineer

Webisoft | Apr 2024 – Present

- Built production-grade LLM systems using RAG, LangChain, and LangGraph, improving response accuracy by 40%.
- Architected AI inference services using FastAPI, supporting high-throughput, low-latency workloads.
- Implemented advanced vector retrieval systems (Pinecone, FAISS) for semantic search and knowledge retrieval.
- Reduced inference latency by 30% through caching, batching, and prompt optimization techniques.
- Designed multi-agent AI systems for workflow automation and decision-making.
- Deployed scalable AI infrastructure on AWS using Docker, Kubernetes, and CI/CD pipelines.
- Automated AI workflows using n8n and event-driven orchestration pipelines.

AI Engineer

Confiz | Nov 2021 – Mar 2024

- Delivered production AI systems using Next.js, Node.js, Python, and TypeScript.
- Built RAG-based conversational systems using LangChain and vector databases.
- Integrated OpenAI and Hugging Face APIs for automation and generative AI features.
- Improved backend performance by 35% via query optimization and caching.
- Designed scalable REST and GraphQL APIs with PostgreSQL and MongoDB.
- Deployed cloud-native applications on AWS with CI/CD pipelines and monitoring.

Full Stack Developer

AcclivousByte | Feb 2021 – Oct 2021

- Built full-stack applications using Python, Node.js, and React.js, improving feature delivery speed and scalability.
- Developed and enhanced REST APIs and backend services, increasing system performance and handling higher request loads.
- Optimized PostgreSQL queries, reducing response time by 25% and improving data retrieval efficiency.
- Integrated third-party APIs and automated data workflows, reducing manual effort and improving processing speed.
- Deployed applications on AWS, achieving high availability (99.9%) and reliable production releases.



Projects

SoundhoundAI – Enterprise Conversational AI Platform

Technologies: Python, FastAPI, RAG, LangGraph, Vector Databases, AWS

- Built enterprise AI assistant for automated workflows and customer support.
- Designed multi-agent orchestration using LangGraph.
- Implemented RAG pipelines for context-aware retrieval and response generation.

Sabermine – AI Document Processing System

Technologies: Python, FastAPI, OCR, NLP, AWS

- Engineered document intelligence system for structured data extraction.
- Processed PDFs and scanned documents using OCR + NLP pipelines.
- Built scalable backend for high-volume document ingestion.

Myaigi – Multi-Agent AI System

Technologies: Node.js, JavaScript, LangChain, LangGraph, LLMs, RAG

- Developed multi-agent AI orchestration system for complex workflows.
- Built modular agent architecture for task-specific AI execution.
- Implemented reasoning pipelines using LLM-based decision systems.

Highline Residential – Real Estate Platform

Technologies: Python, FastAPI, JavaScript, React.js, PostgreSQL

- Built backend APIs for property data processing and search systems.
- Developed React dashboards for real estate listings and analytics.
- Improved PostgreSQL query performance for faster data retrieval (30% faster).

Certifications

- ChatGPT Prompt Engineering – DeepLearning.AI
- LangChain & LLM Application Development
- AWS Certified Developer – Associate

EDUCATION

Bachelor's degree (Computer Science)

University of Engineering and Technology (UET), Lahore, Pakistan

